

# Convolutional Neural Network

(Draft)

Quang-Vinh Dinh  
Ph.D. in Computer Science

# MLP for Fashion-MNIST



## ❖ Sigmoid and SGD

```
# model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation='sigmoid',
                        kernel_initializer=keras.initializers.RandomNormal(stddev=0.01),
                        bias_initializer=keras.initializers.Zeros()),
    keras.layers.Dense(10, activation='softmax',
                        kernel_initializer=keras.initializers.RandomNormal(stddev=0.01),
                        bias_initializer=keras.initializers.Zeros())
])

model.compile(optimizer='sgd',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
history_data = model.fit(train_images, train_labels,
                          validation_data=(test_images, test_labels),
                          batch_size=1024, epochs=80, verbose=2)
```

# MLP for Fashion-MNIST

## ❖ Sigmoid and SGD



Accuracy: 0.7075 - Val\_accuracy: 0.6988

# MLP for Fashion-MNIST

## ❖ Xavier, Sigmoid and SGD



```
# model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='sgd',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history_data = model.fit(train_images, train_labels,
                        validation_data=(test_images, test_labels),
                        batch_size=1024, epochs=50, verbose=2)
```

# MLP for Fashion-MNIST

## ❖ Xavier, Sigmoid and SGD



Accuracy: 0.7578 - Val\_accuracy: 0.7480

# MLP for Fashion-MNIST

## ❖ Sigmoid and Adam



```
# model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history_data = model.fit(train_images, train_labels,
                        validation_data=(test_images, test_labels),
                        batch_size=1024, epochs=50, verbose=2)
```

# MLP for Fashion-MNIST

## ❖ Sigmoid and Adam



Accuracy: 0.9185 ; Val\_accuracy: 0.8836

# MLP for Fashion-MNIST

## ❖ Sigmoid and Adam: More layers

```
# model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(10, activation='softmax')
])



model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history_data = model.fit(train_images, train_labels,
                        validation_data=(test_images, test_labels),
                        batch_size=1024, epochs=50, verbose=2)
```



# MLP for Fashion-MNIST

## ❖ Sigmoid and Adam: More layers



Accuracy: 0.9230 - Val\_accuracy: 0.8856

# MLP for Fashion-MNIST

## ❖ Sigmoid and Adam: Keep adding more layers



```
# model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history_data = model.fit(train_images, train_labels,
                        validation_data=(test_images, test_labels),
                        batch_size=1024, epochs=50, verbose=2)
```

# MLP for Fashion-MNIST

❖ Sigmoid and Adam: Keep adding more layers



Accuracy: 0.9202 - Val\_accuracy: 0.8686

# MLP for Fashion-MNIST

## ❖ Sigmoid and Adam: Keep adding more layers



```
# model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(128, activation='sigmoid'),
    keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history_data = model.fit(train_images, train_labels,
                        validation_data=(test_images, test_labels),
                        batch_size=1024, epochs=250, verbose=0)
```

# MLP for Fashion-MNIST

❖ Sigmoid and Adam: Keep adding more layers



# MLP for Fashion-MNIST



- ❖ ReLU and Adam
- ❖ One hidden layer

```
# model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
history_data = model.fit(train_images, train_labels,
                        validation_data=(test_images, test_labels),
                        batch_size=1024, epochs=150, verbose=2)
```

# MLP for Fashion-MNIST

- ❖ ReLU and Adam
- ❖ One hidden layer



Accuracy: 0.9735 - Val\_accuracy: 0.8873

# MLP for Fashion-MNIST

- ❖ ReLU and Adam
- ❖ Three hidden layers

```
# model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(10, activation='softmax')
])



model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history_data = model.fit(train_images, train_labels,
                        validation_data=(test_images, test_labels),
                        batch_size=1024, epochs=150, verbose=2)
```



# MLP for Fashion-MNIST

- ❖ ReLU and Adam
- ❖ Three hidden layers



Accuracy: 0.9914 - Val\_accuracy: 0.8858

# MLP for Cifar10



- ❖ ReLU and Adam
- ❖ One hidden layer

```
# model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(32, 32, 3)),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
history_data = model.fit(train_images, train_labels,
                        validation_data=(test_images, test_labels),
                        batch_size=1024, epochs=150, verbose=2)
```

# MLP for Cifar10

- ❖ ReLU and Adam
- ❖ One hidden layer



Accuracy: 0.5420 - Val\_accuracy: 0.4748

# MLP for Cifar10

- ❖ ReLU and Adam
- ❖ Two hidden layers



```
# model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(32, 32, 3)),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history_data = model.fit(train_images, train_labels,
                        validation_data=(test_images, test_labels),
                        batch_size=1024, epochs=150, verbose=2)
```

# MLP for Cifar10

- ❖ ReLU and Adam
- ❖ Two hidden layers



Accuracy: 0.6374 - Val\_accuracy: 0.5144

# MLP for Cifar10

❖ ReLU and Adam

❖ Five hidden layers



```
# model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(32, 32, 3)),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(128, activation='relu'),
    keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history_data = model.fit(train_images, train_labels,
                        validation_data=(test_images, test_labels),
                        batch_size=1024, epochs=150, verbose=2)
```

# MLP for Cifar10

- ❖ ReLU and Adam
- ❖ Five hidden layers



Accuracy: 0.8560 - Val\_accuracy: 0.4845


# Outline

- **From MLP to CNN**
- **Feature Map Down-sampling**
- **Padding**
- **1x1 Convolution**
- **Image classification: Cifar-10 data**



# From MLP to CNN



## ❖ Multi-layer Perceptron



Problem: Remove spatial information of the data  
Inefficiently have a large amount of parameters

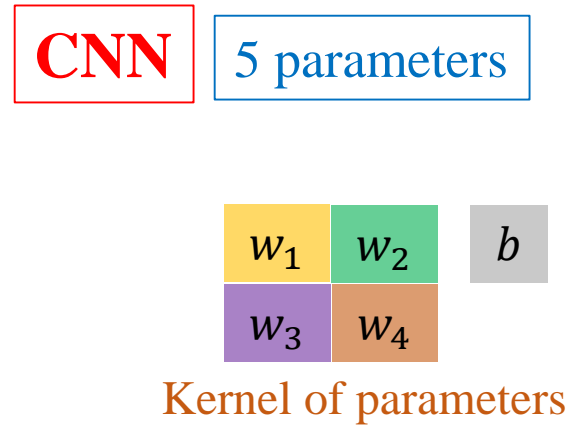
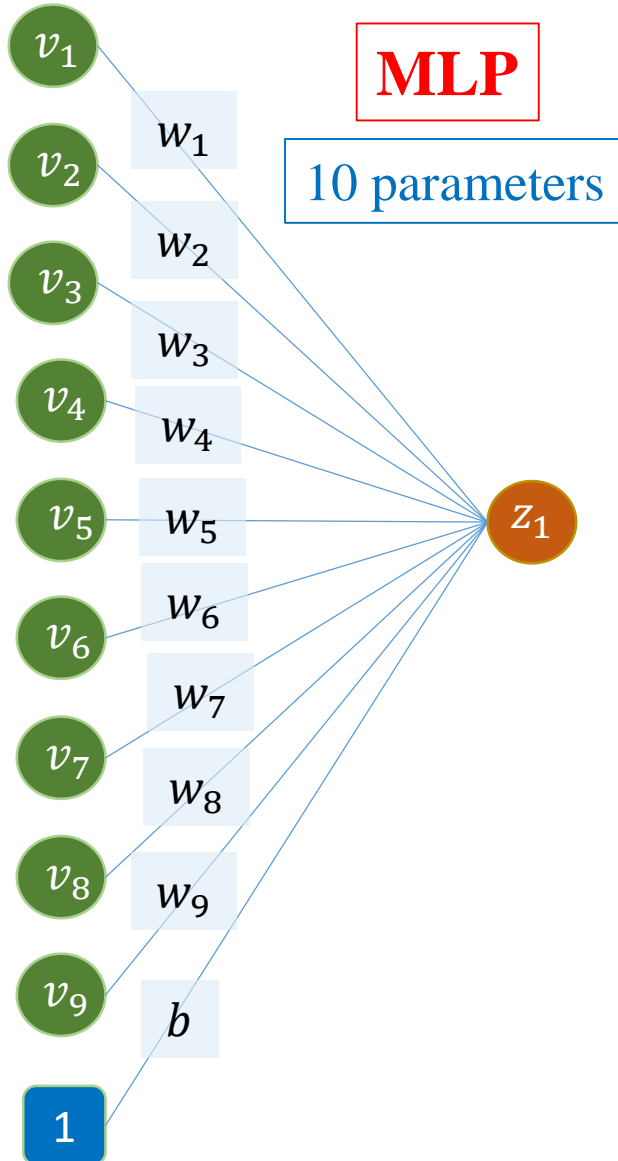
# From MLP to CNN

## ❖ Problem of flattening data





Remove spatial information of the data

# From MLP to CNN

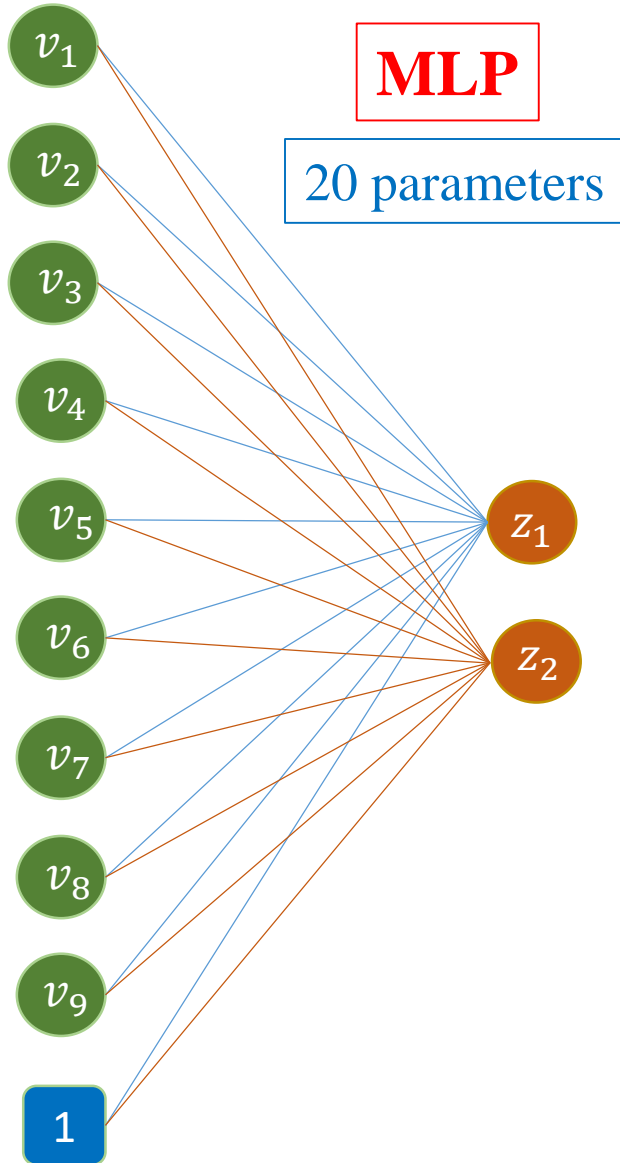


$$m_1 = v_1 w_1 + v_2 w_2 + v_4 w_3 + v_5 w_4 + b$$








Feature Map

# From MLP to CNN




**CNN** 10 parameters



Kernel 1  $\neq$  Kernel 2

# From MLP to CNN

## ❖ Understand convolution



(Height=3, Width=3, Channel=1)

Shape=(3,3,1)

\*





Shape=(2,2,1)

#parameters (including bias) = 5

#channels of data  $\overset{\text{must}}{=}$  #channels of kernel

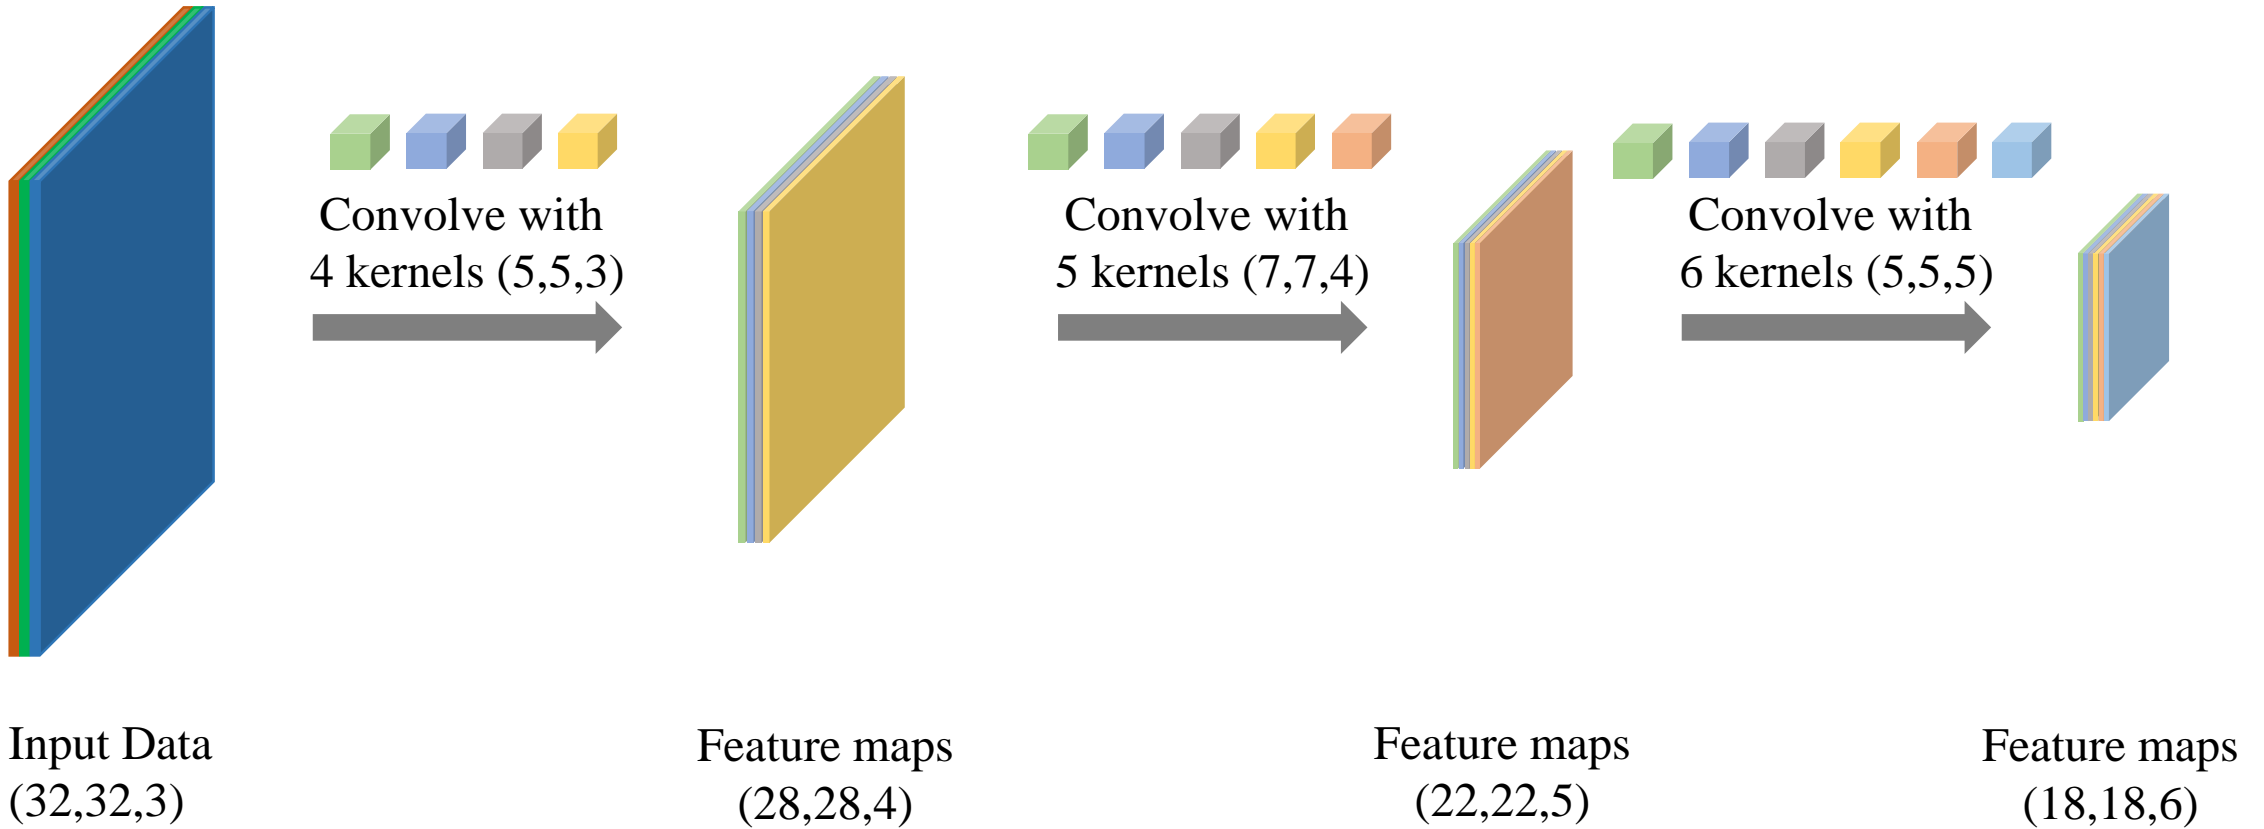
# From MLP to CNN

## ❖ Understand convolution



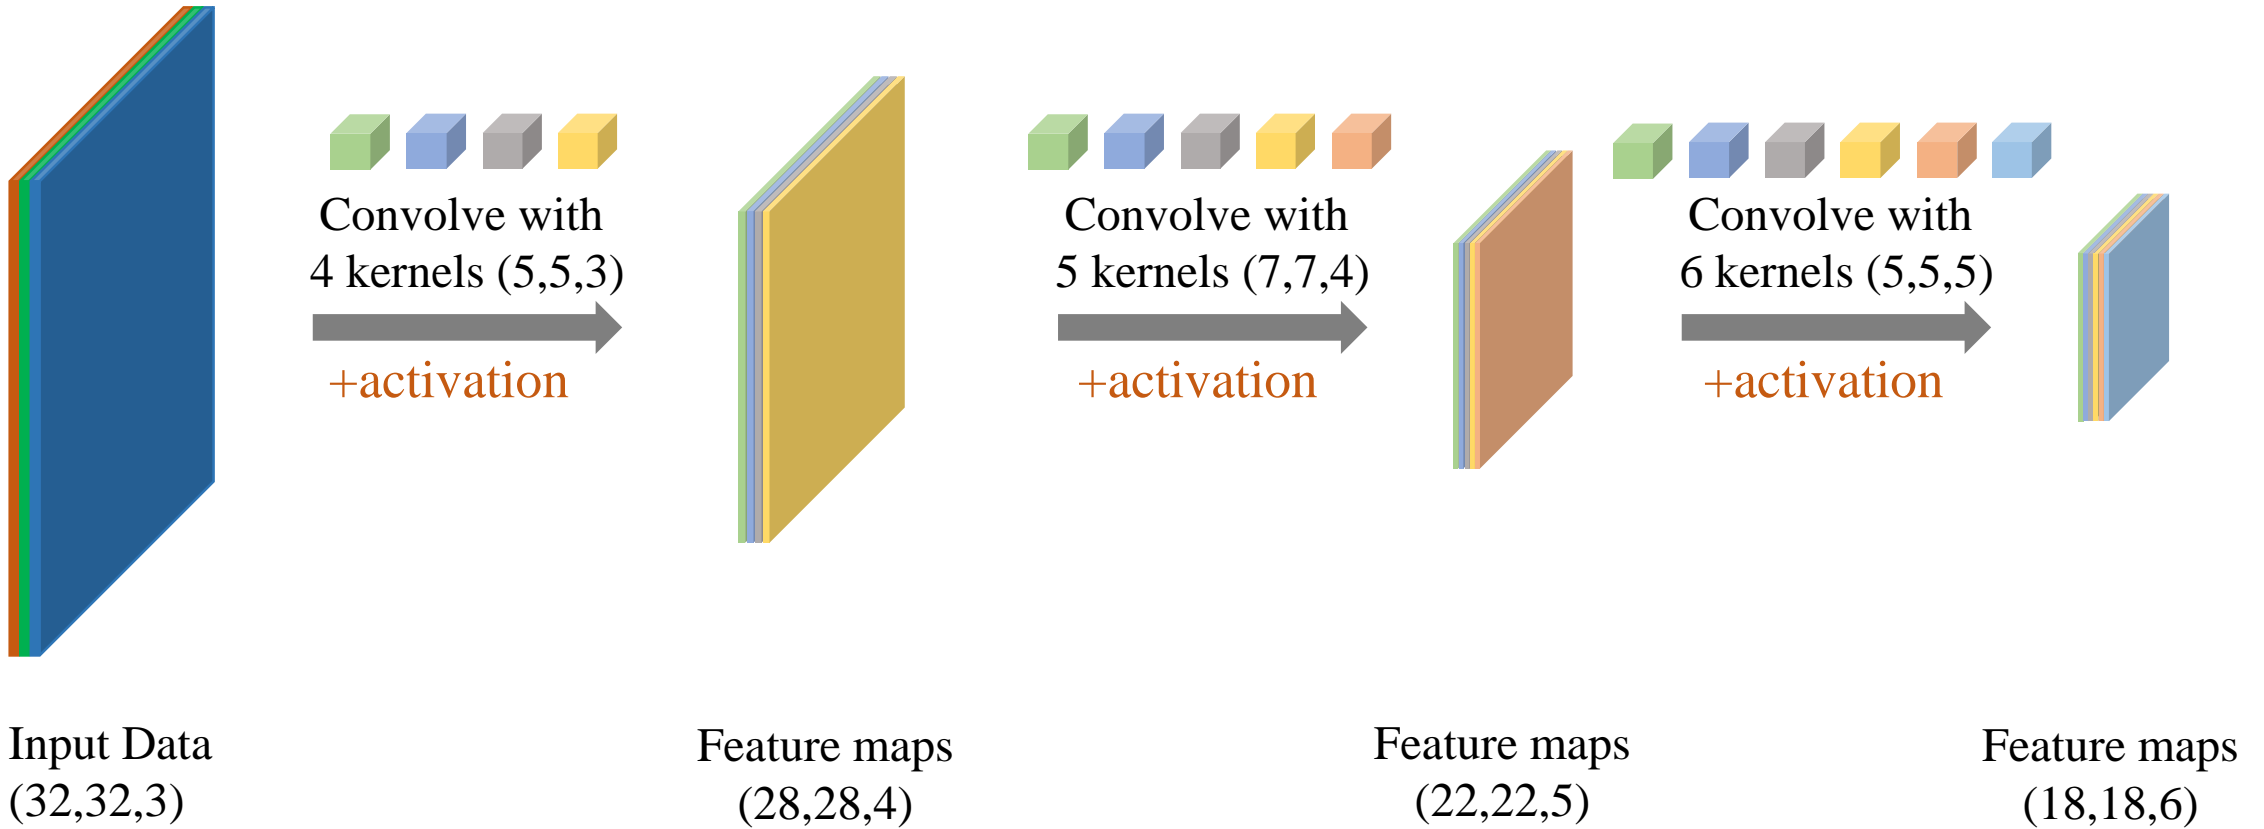
# From MLP to CNN

## ❖ A stack of convolutions



# From MLP to CNN

## ❖ A stack of pairs of convolution+activation







# From MLP to CNN


## ❖ Convolution layer in Keras

```
keras.layers.Conv2D(filters=1, kernel_size=5, activation='relu')
```



Input Data  
(32,32,3)


  
Convolve with  
1 kernel (5,5,1)  
→  
+activation



Feature map  
(28,28,1)



Input Data  
(32,32,3)

  
Convolve with  
1 kernel (5,5,3)  
→  
+activation




Feature map  
(28,28,1)


# From MLP to CNN

## ❖ Convolution layer in Keras

```
keras.layers.Conv2D(filters=4, kernel_size=5, activation='relu')
```




Input Data  
(32,32,3)



Feature maps  
(28,28,4)



Input Data  
(32,32,3)



Feature maps  
(28,28,4)

# From MLP to CNN

## Fashion-MNIST dataset

Grayscale images

Resolution=28x28

Training set: 60000 samples

Testing set: 10000 samples


T-shirt



Trouser



Pullover



Dress




Coat



Sandal



Shirt



Sneaker



Bag



Ankle  
Boot



# From MLP to CNN

```
1 # model
2 model = keras.models.Sequential()
3
4 # input is with the shape of (28, 28, 1)
5 model.add(tf.keras.Input(shape=(28, 28, 1)))
6
7 # Convolve with 32 (7x7) kernel; Output: (22x22x32)
8 model.add(keras.layers.Conv2D(32, (7, 7), activation='relu'))
9 # Convolve with 64 (7x7) kernel; Output: (16x16x64)
10 model.add(keras.layers.Conv2D(64, (7, 7), activation='relu'))
11 # Convolve with 128 (7x7) kernel; Output: (10x10x128)
12 model.add(keras.layers.Conv2D(128, (7, 7), activation='relu'))
13 # Convolve with 256 (7x7) kernel; Output: (4x4x256)
14 model.add(keras.layers.Conv2D(256, (7, 7), activation='relu'))
15
16 # flatten
17 model.add(keras.layers.Flatten())
18 model.add(keras.layers.Dense(10, activation='softmax'))
19
20 # compile and train
21 model.compile(optimizer='adam', metrics=['accuracy'],
22               loss='sparse_categorical_crossentropy')
23 model.fit(train_images, train_labels, epochs=10)
24
25 # testing
26 test_loss, test_acc = model.evaluate(test_images,
27                                     test_labels, verbose=2)
28 print('Test accuracy:', test_acc)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 22, 22, 32)	1600
conv2d_1 (Conv2D)	(None, 16, 16, 64)	100416
conv2d_2 (Conv2D)	(None, 10, 10, 128)	401536
conv2d_3 (Conv2D)	(None, 4, 4, 256)	1605888
flatten (Flatten)	(None, 4096)	0
dense (Dense)	(None, 10)	40970

Total params: 2,150,410

Trainable params: 2,150,410

Non-trainable params: 0

Train on 60000 samples

Epoch 1/10

60000/60000 - 577s 10ms/sample - loss: 0.5021 - accuracy: 0.8138

Epoch 2/10

60000/60000 - 578s 10ms/sample - loss: 0.3388 - accuracy: 0.8757

Epoch 3/10

60000/60000 - 567s 9ms/sample - loss: 0.2993 - accuracy: 0.8880

Epoch 4/10

60000/60000 - 545s 9ms/sample - loss: 0.2726 - accuracy: 0.8995

Epoch 5/10

60000/60000 - 1254s 21ms/sample - loss: 0.2475 - accuracy: 0.9083

Epoch 6/10

60000/60000 - 563s 9ms/sample - loss: 0.2201 - accuracy: 0.9172

Epoch 7/10

60000/60000 - 571s 10ms/sample - loss: 0.1983 - accuracy: 0.9254

Epoch 8/10

60000/60000 - 581s 10ms/sample - loss: 0.1806 - accuracy: 0.9340

Epoch 9/10

60000/60000 - 581s 10ms/sample - loss: 0.1517 - accuracy: 0.9431

Epoch 10/10

60000/60000 - 2145s 36ms/sample - loss: 0.1378 - accuracy: 0.9495

10000/1 - 19s - loss: 1.2228 - accuracy: 0.8858

Test accuracy: 0.8858

# Outline

- **From MLP to CNN**
- **Feature Map Down-sampling**
- **Padding**
- **1x1 Convolution**
- **Image classification: Cifar-10 data**

# Down-sample Feature Map

## Max pooling: Features are preserved

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

Data

2x2 max  
pooling

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

) =

$m_1$	$m_2$
$m_3$	$m_4$

$$m_1 = \max(v_1, v_2, v_5, v_6)$$

$$m_2 = \max(v_3, v_4, v_7, v_8)$$

$$m_3 = \max(v_9, v_{10}, v_{13}, v_{14})$$

$$m_4 = \max(v_{11}, v_{12}, v_{15}, v_{16})$$



Feature map (220x220)

max pooling  
(2x2)



Feature map  
(110x110)

max pooling  
(2x2)



Feature map  
(55x55)

# Down-sample Feature Map

## Max pooling: Features are preserved

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

Data

2x2 max  
pooling

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

) =

$m_1$	$m_2$
$m_3$	$m_4$

$$m_1 = \max(v_1, v_2, v_5, v_6)$$

$$m_2 = \max(v_3, v_4, v_7, v_8)$$

$$m_3 = \max(v_9, v_{10}, v_{13}, v_{14})$$

$$m_4 = \max(v_{11}, v_{12}, v_{15}, v_{16})$$

`keras.layers.MaxPooling2D(pool_size=2)`



Feature map (220x220)

max pooling  
(2x2)



Feature map  
(110x110)

max pooling  
(2x2)




Feature map  
(55x55)



# Down-sample Feature Map

## Example 1: Using max pooling





# Down-sample Feature Map

## ❖ Example 1: Using max pooling

```
Epoch 1/10
60000/60000 [=====] - 281s 5ms/sample - loss: 0.4243 - accuracy: 0.8423
Epoch 2/10
60000/60000 [=====] - 295s 5ms/sample - loss: 0.2678 - accuracy: 0.9029
Epoch 3/10
60000/60000 [=====] - 305s 5ms/sample - loss: 0.2200 - accuracy: 0.9199
Epoch 4/10
60000/60000 [=====] - 298s 5ms/sample - loss: 0.1832 - accuracy: 0.9330
Epoch 5/10
60000/60000 [=====] - 298s 5ms/sample - loss: 0.1510 - accuracy: 0.9447
Epoch 6/10
60000/60000 [=====] - 302s 5ms/sample - loss: 0.1245 - accuracy: 0.9540
Epoch 7/10
60000/60000 [=====] - 306s 5ms/sample - loss: 0.0988 - accuracy: 0.9639
Epoch 8/10
60000/60000 [=====] - 292s 5ms/sample - loss: 0.0847 - accuracy: 0.9684
Epoch 9/10
60000/60000 [=====] - 291s 5ms/sample - loss: 0.0709 - accuracy: 0.9741
Epoch 10/10
60000/60000 [=====] - 289s 5ms/sample - loss: 0.0606 - accuracy: 0.9786
10000/1 - 8s - loss: 0.4003 - accuracy: 0.9125
```

Test accuracy: 0.9125

# Down-sample Feature Map

❖ Convolve with stride

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

Data **D**

$w_1$	$w_2$	$b$
$w_3$	$w_4$	

Kernel of parameters

Convolve **D**  
with stride=1

$m_1$	$m_2$	$m_3$
$m_4$	$m_5$	$m_6$
$m_7$	$m_8$	$m_9$

Output

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

# Down-sample Feature Map

## ❖ Convolve with stride

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

Data D

$w_1$	$w_2$	$b$
$w_3$	$w_4$	

Kernel of parameters

Convolve D  
with stride=2

$m_1$	$m_2$
$m_3$	$m_4$

Output

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

```
keras.layers.Conv2D(32, (3, 3), strides=1, activation='relu')
```

```
keras.layers.Conv2D(32, (3, 3), strides=2, activation='relu')
```


# Down-sample Feature Map

## Example 2: Using Conv with strides

```
# model
model = keras.models.Sequential()
model.add(tf.keras.Input(shape=(28, 28, 1)))

model.add(keras.layers.Conv2D(32, (3, 3), strides=1, activation='relu'))
model.add(keras.layers.Conv2D(64, (3, 3), strides=2, activation='relu'))
model.add(keras.layers.Conv2D(128, (3, 3), strides=1, activation='relu'))
model.add(keras.layers.Conv2D(256, (3, 3), strides=2, activation='relu'))
model.add(keras.layers.Conv2D(512, (3, 3), strides=1, activation='relu'))

# flatten
model.add(keras.layers.Flatten())
model.add(keras.layers.Dense(10, activation='softmax'))
```



# Down-sample Feature Map

## ❖ Example 2: Using Conv with strides

```
Epoch 1/10
60000/60000 [=====] - 276s 5ms/sample - loss: 0.4500 - accuracy: 0.8350
Epoch 2/10
60000/60000 [=====] - 268s 4ms/sample - loss: 0.2934 - accuracy: 0.8922
Epoch 3/10
60000/60000 [=====] - 285s 5ms/sample - loss: 0.2376 - accuracy: 0.9108
Epoch 4/10
60000/60000 [=====] - 291s 5ms/sample - loss: 0.1942 - accuracy: 0.9283
Epoch 5/10
60000/60000 [=====] - 288s 5ms/sample - loss: 0.1536 - accuracy: 0.9420
Epoch 6/10
60000/60000 [=====] - 288s 5ms/sample - loss: 0.1198 - accuracy: 0.9548
Epoch 7/10
60000/60000 [=====] - 289s 5ms/sample - loss: 0.0934 - accuracy: 0.9653
Epoch 8/10
60000/60000 [=====] - 199s 3ms/sample - loss: 0.0746 - accuracy: 0.9730
Epoch 9/10
60000/60000 [=====] - 200s 3ms/sample - loss: 0.0625 - accuracy: 0.9772
Epoch 10/10
60000/60000 [=====] - 199s 3ms/sample - loss: 0.0566 - accuracy: 0.9797
10000/1 - 5s - loss: 0.3381 - accuracy: 0.9055
```

Test accuracy: 0.9055

# Outline

- **From MLP to CNN**
- **Feature Map Down-sampling**
- **Padding**
- **1x1 Convolution**
- **Image classification: Cifar-10 data**

# Padding

**Goal: Keep resolution  
of feature map**

$v_1$	$v_2$	$v_3$	$v_4$
$v_5$	$v_6$	$v_7$	$v_8$
$v_9$	$v_{10}$	$v_{11}$	$v_{12}$
$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$

Data  $D$   
(4x4)

$w_1$	$w_2$	$w_3$
$w_4$	$w_5$	$w_6$
$w_7$	$w_8$	$w_9$

Kernel of parameters

$b$

Without using padding  
or padding=0

Convolve  
with stride=1 ( $D$ ) =

$m_1$	$m_2$
$m_4$	$m_5$

Output  
(2x2)

Padding = 1

$v$	$v$	$v$	$v$	$v$	$v$
$v$	$v_1$	$v_2$	$v_3$	$v_4$	$v$
$v$	$v_5$	$v_6$	$v_7$	$v_8$	$v$
$v$	$v_9$	$v_{10}$	$v_{11}$	$v_{12}$	$v$
$v$	$v_{13}$	$v_{14}$	$v_{15}$	$v_{16}$	$v$
$v$	$v$	$v$	$v$	$v$	$v$

Data  $D_p$

Convolve  
with stride=1 ( $D_p$ ) =

$m_1$	$m_2$	$m_3$	$m_4$
$m_5$	$m_6$	$m_7$	$m_8$
$m_9$	$m_{10}$	$m_{11}$	$m_{12}$
$m_{13}$	$m_{14}$	$m_{15}$	$m_{16}$

Output  
(4x4)

# Padding

## Example

```
# model
model = keras.models.Sequential()
model.add(tf.keras.Input(shape=(28, 28, 1)))

model.add(keras.layers.Conv2D(32, (3, 3), strides=1, padding='same', activation='relu'))
model.add(keras.layers.Conv2D(64, (3, 3), strides=2, padding='same', activation='relu'))


model.add(keras.layers.Conv2D(128, (3, 3), strides=1, padding='same', activation='relu'))
model.add(keras.layers.Conv2D(256, (3, 3), strides=2, padding='same', activation='relu'))

model.add(keras.layers.Conv2D(512, (3, 3), strides=1, padding='same', activation='relu'))

# flatten
model.add(keras.layers.Flatten())
model.add(keras.layers.Dense(10, activation='softmax'))
```

(3x3) Conv with stride=1,  
padding='same' + ReLU

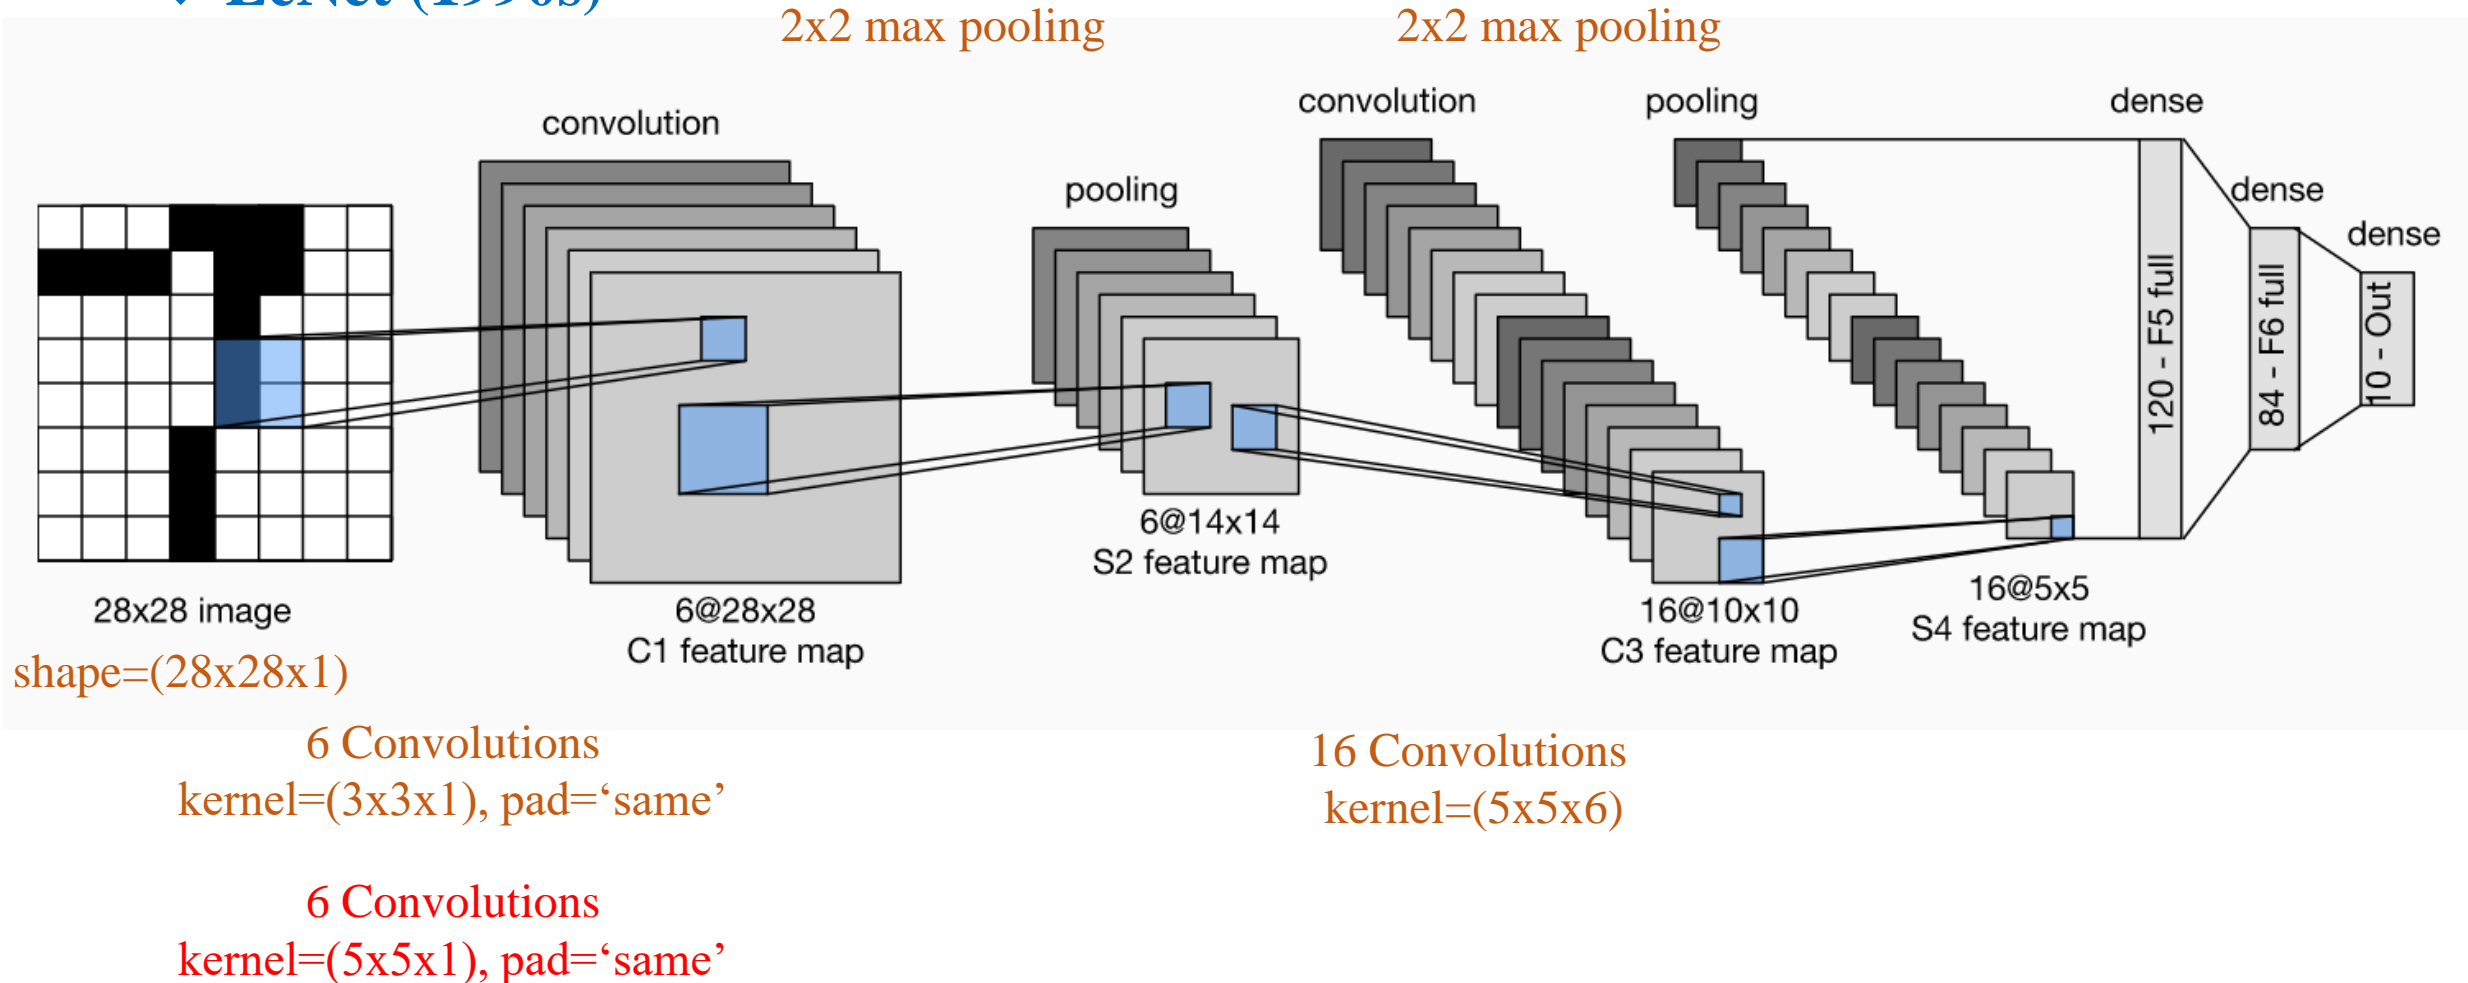
(3x3) Conv with stride=2,  
padding='same' + ReLU





# Padding


## ❖ LeNet (1990s)



# Padding

## ❖ LeNet (1990s)

[https://d2l.ai/chapter\\_convolutional-neural-networks/lenet.html](https://d2l.ai/chapter_convolutional-neural-networks/lenet.html)



```
# model architecture
model = tf.keras.Sequential()
# input shape (28,28,1)
model.add(tf.keras.Input(shape=(28, 28, 1)))

# convolution 1 and max pooling 1
model.add(tf.keras.layers.Conv2D(6, (5,5), padding='same', activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

# convolution 2 and max pooling 2
model.add(tf.keras.layers.Conv2D(filters=16, kernel_size=5, activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

# Flatten
model.add(tf.keras.layers.Flatten())

# fully connected
model.add(tf.keras.layers.Dense(120, activation='relu'))
model.add(tf.keras.layers.Dense(84, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation='softmax'))
```

# Outline

- **From MLP to CNN**
- **Feature Map Down-sampling**
- **Padding**
- **1x1 Convolution**
- **Image classification: Cifar-10 data**

# 1x1 Convolution

## ❖ Why 1x1 Convolution

### ❖ Flexible input size



Yann LeCun

April 7, 2015 · 🌐



In Convolutional Nets, there is no such thing as "fully-connected layers". There are only convolution layers with 1x1 convolution kernels and a full connection table.

It's a too-rarely-understood fact that ConvNets don't need to have a fixed-size input. You can train them on inputs that happen to produce a single output vector (with no spatial extent), and then apply them to larger images. Instead of a single output vector, you then get a spatial map of output vectors. Each vector sees input windows at different locations on the input. In that scenario, the "fully connected layers" really act as 1x1 convolutions.

Yann LeCun



Yann LeCun in 2018

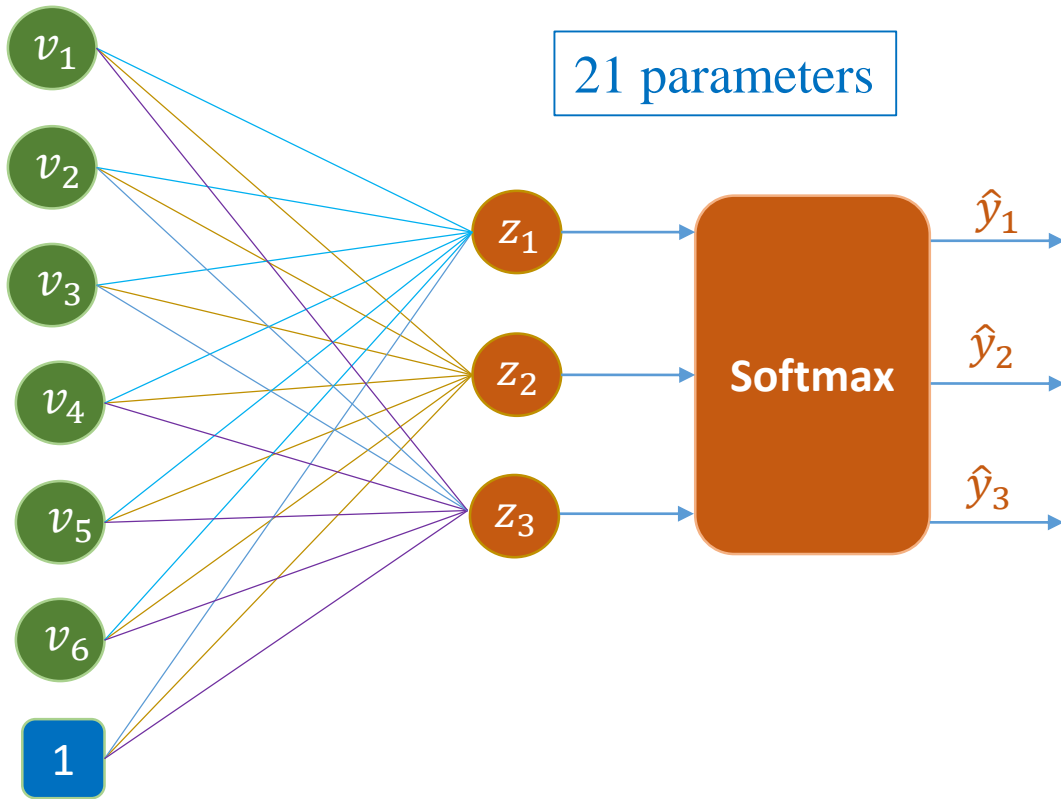
Born	July 8, 1960 (age 60) <a href="#">Soisy-sous-Montmorency, France</a>
Alma mater	<a href="#">ESIEE Paris</a> (MSc) <a href="#">Pierre and Marie Curie University</a> (PhD)
Known for	<a href="#">Deep learning</a>
Awards	<a href="#">Turing Award</a> (2018) <a href="#">AAAI Fellow</a> (2019) <a href="#">Legion of Honour</a> (2020)
<b>Scientific career</b>	
Institutions	<a href="#">Bell Labs</a> (1988-1996) <a href="#">New York University</a> <a href="#">Facebook</a>
Thesis	<i>Modèles connexionnistes de l'apprentissage (connectionist learning models)</i> (1987a)
Doctoral advisor	<a href="#">Maurice Milgram</a>
Website	<a href="http://yann.lecun.com">yann.lecun.com</a>

# 1x1 Convolution

## ❖ Comparison

### Fully connected layer

21 parameters



### CNN

21 parameters



# 1x1 Convolution


## Replace FC by 1x1 Conv

```
# model
model = keras.models.Sequential()
model.add(tf.keras.Input(shape=(28, 28, 1)))
```

```
model.add(keras.layers.Conv2D(32, (3, 3), strides=1, padding='same', activation='relu'))
model.add(keras.layers.Conv2D(64, (3, 3), strides=2, padding='same', activation='relu'))
```

```
model.add(keras.layers.Conv2D(128, (3, 3), strides=1, padding='same', activation='relu'))
model.add(keras.layers.Conv2D(256, (3, 3), strides=2, padding='same', activation='relu'))
```

```
model.add(keras.layers.Conv2D(512, (7, 7), strides=1, activation='relu'))
model.add(keras.layers.Conv2D(10, (1, 1), strides=1, activation='softmax'))
```



# 1x1 Convolution

## Dynamic input sizes

```
# model
model = keras.models.Sequential()
model.add(tf.keras.Input(shape=(None, None, 1)))

model.add(keras.layers.Conv2D(32, (3, 3), strides=1, padding='same', activation='relu'))
model.add(keras.layers.Conv2D(64, (3, 3), strides=2, padding='same', activation='relu'))

model.add(keras.layers.Conv2D(128, (3, 3), strides=1, padding='same', activation='relu'))
model.add(keras.layers.Conv2D(256, (3, 3), strides=2, padding='same', activation='relu'))

model.add(keras.layers.Conv2D(512, (7, 7), strides=1, activation='relu'))
model.add(keras.layers.Conv2D(10, (1, 1), strides=1, activation='softmax'))
```

```
1 shape = (32,28,28,1)
2 data = tf.random.normal(shape)
3 print(data.shape)
4
5 output = model.predict(data)
6 print(output.shape)
```

```
(32, 28, 28, 1)
(32, 1, 1, 10)
```

```
1 shape = (32,64,64,1)
2 data = tf.random.normal(shape)
3 print(data.shape)
4
5 output = model.predict(data)
6 print(output.shape)
```

```
(32, 64, 64, 1)
(32, 10, 10, 10)
```

Shape=(batch size, height, width, channel)



# 1x1 Convolution

## Dynamic input sizes

```
# model
model = keras.models.Sequential()
model.add(tf.keras.Input(shape=(None, None, 1)))

model.add(keras.layers.Conv2D(32, (3, 3), activation='relu'))
model.add(keras.layers.Conv2D(64, (3, 3), activation='relu'))
model.add(keras.layers.MaxPooling2D(2))

model.add(keras.layers.Conv2D(128, (3, 3), activation='relu'))
model.add(keras.layers.Conv2D(256, (3, 3), activation='relu'))
model.add(keras.layers.MaxPooling2D(2))

model.add(keras.layers.Conv2D(512, (3, 3), activation='relu'))

# flatten
model.add(keras.layers.Flatten())
model.add(keras.layers.Dense(10, activation='softmax'))
```

```
C:\installed\lib\site-packages\tensorflow_core\python\keras\layers\core.py in build(self, input_shape)
    1013     input_shape = tensor_shape.TensorShape(input_shape)
    1014     if tensor_shape.dimension_value(input_shape[-1]) is None:
-> 1015         raise ValueError('The last dimension of the inputs to `Dense` '
    1016                           'should be defined. Found `None`.')
    1017     last_dim = tensor_shape.dimension_value(input_shape[-1])
```

**ValueError:** The last dimension of the inputs to `Dense` should be defined. Found `None`.



# Outline

- **From MLP to CNN**
- **Feature Map Down-sampling**
- **Padding**
- **1x1 Convolution**
- **Image classification: Cifar-10 data**

# Image Classification

Cifar-10 dataset

Color images  
Resolution=32x32  
Training set: 50000 samples  
Testing set: 10000 samples

airplane



automobile



bird



cat



deer



dog



frog



horse




ship



truck



# Cifar-10 Image Classification



# Image Classification

## Cifar-10

```
1 import tensorflow as tf
2
3 # data preparation
4 cifar10 = tf.keras.datasets.cifar10
5 (x_train, y_train), (x_test, y_test) = cifar10.load_data()
6
7 # normalize
8 x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
# model
model = keras.models.Sequential()
model.add(tf.keras.Input(shape=(height, width, 3)))

model.add(keras.layers.Conv2D(32, (3, 3), strides=1, padding='same', activation='relu'))
model.add(keras.layers.Conv2D(64, (3, 3), strides=2, padding='same', activation='relu'))

model.add(keras.layers.Conv2D(128, (3, 3), strides=1, padding='same', activation='relu'))
model.add(keras.layers.Conv2D(256, (3, 3), strides=2, padding='same', activation='relu'))

model.add(keras.layers.Conv2D(512, (3, 3), strides=1, padding='same', activation='relu'))
model.add(keras.layers.Conv2D(1024, (3, 3), strides=2, padding='same', activation='relu'))

# flatten
model.add(keras.layers.Flatten())
model.add(keras.layers.Dense(512, activation='relu'))
model.add(keras.layers.Dense(10, activation='softmax'))
model.summary()
```

# Image Classification

## ❖ Demo

```
Python 3.7.3 (default, Apr 24 2019, 15:29:51) [MSC v.1915 64 bit (AMD64)] ::  
Type "help", "copyright", "credits" or "license" for more information.  
>>>  
>>>  
>>>  
>>>  
>>>  
>>>  
>>>  
>>>  
>>> for epoch in range(n_epochs):  
...     sum_of_losses = 0  
...     gradients = np.zeros((2,1))  
...  
...     for index in range(4):  
...         xi = X_b[index:index+1]  
...         yi = y[index:index+1]
```

# Reading and Exercises

## ❖ Exercises

- 1) Use LeNet for the fashion-MNIST and Cifar-10 data sets
- 2) What are the advantages of using 1x1 Conv instead of FC

## ❖ Reading

<https://cs231n.github.io/convolutional-networks/>

<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

